

```
In [1]: #####  
## Title : 'Regression Analysis'  
## Author : 'Bushra Tariq Kiyani'  
## Date : '27th January, 2023'  
## Output : pdf document  
## Group : 5  
#####
```

```
In [3]: library(ggplot2)  
library(GGally)  
library(ggpubr)  
library(dplyr)  
library(olsrr)  
library(reshape2)
```

```
In [4]: #Loading the Input data  
input_data <- read.csv('bodymeasurements.csv')
```

```
In [5]: #Remove the variable ID  
input_data <- input_data %>% dplyr::select(-ID)
```

```
In [6]: #View the input data  
head(input_data)
```

A data.frame: 6 × 12

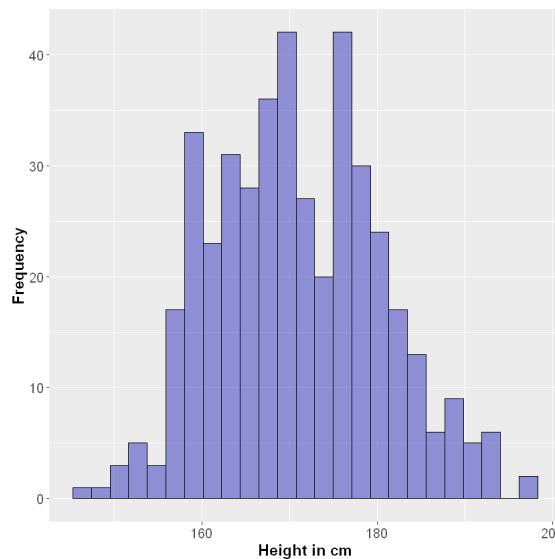
| | Age | Sex | Height | Chest | Belly | Thigh | Knee | Calf | Ankle | Biceps | Wrist | Weight |
|---|-------|-------|--------|-------|-------|-------|-------|-------|-------|--------|-------|--------|
| | <int> | <chr> | <dbl> | <dbl> | <dbl> | <dbl> | <dbl> | <dbl> | <dbl> | <dbl> | <dbl> | <dbl> |
| 1 | 21 | m | 174.0 | 89.5 | 71.5 | 51.5 | 34.5 | 36.5 | 23.5 | 32.5 | 16.5 | 65.6 |
| 2 | 23 | m | 175.3 | 97.0 | 79.0 | 51.5 | 36.5 | 37.5 | 24.5 | 34.4 | 17.0 | 71.8 |
| 3 | 28 | m | 193.5 | 97.5 | 83.2 | 57.3 | 37.0 | 37.3 | 21.9 | 33.4 | 16.9 | 80.7 |
| 4 | 23 | m | 186.5 | 97.0 | 77.8 | 53.0 | 37.0 | 34.8 | 23.0 | 31.0 | 16.6 | 72.6 |
| 5 | 22 | m | 187.2 | 97.5 | 80.0 | 55.4 | 37.7 | 38.6 | 24.4 | 32.0 | 18.0 | 78.8 |
| 6 | 21 | m | 181.5 | 99.9 | 82.5 | 57.5 | 36.6 | 36.1 | 23.5 | 33.0 | 16.9 | 74.8 |

```
In [7]: #Summarise all continous variables  
standard_deviations <- rbind.data.frame(lapply(input_data %>% dplyr::select(-Sex), sd))  
summary_table <- rbind.data.frame(lapply(input_data %>% dplyr::select(-Sex), summary))  
summary_table <- rbind(summary_table, standard_deviations)  
row.names(summary_table) <- c("Min", "Q1", "Median", "Mean", "Q3", "Max", "SD")  
summary_table <- t(round(summary_table, 3))  
summary_table
```

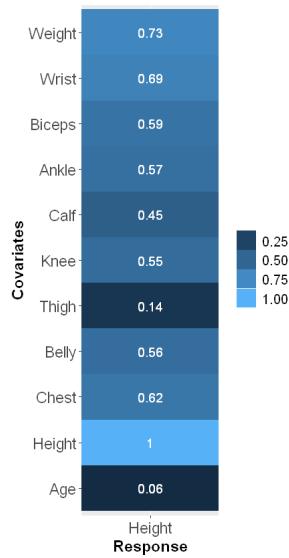
A matrix: 11 × 7 of type dbl

| | Min | Q1 | Median | Mean | Q3 | Max | SD |
|--------|-------|---------|--------|---------|---------|-------|--------|
| Age | 18.0 | 22.000 | 25.00 | 26.896 | 31.250 | 40.0 | 5.849 |
| Height | 147.2 | 163.200 | 170.20 | 170.880 | 177.800 | 198.1 | 9.398 |
| Chest | 72.6 | 84.775 | 90.95 | 92.221 | 99.825 | 116.7 | 9.600 |
| Belly | 57.9 | 67.500 | 74.10 | 75.283 | 82.000 | 113.2 | 9.904 |
| Thigh | 46.3 | 53.700 | 56.30 | 56.822 | 59.500 | 75.7 | 4.431 |
| Knee | 29.0 | 34.300 | 35.90 | 36.005 | 37.700 | 45.7 | 2.544 |
| Calf | 28.4 | 34.000 | 35.80 | 35.880 | 37.700 | 45.0 | 2.756 |
| Ankle | 16.4 | 20.900 | 21.90 | 22.025 | 23.100 | 29.3 | 1.847 |
| Biceps | 22.4 | 27.275 | 30.35 | 30.828 | 34.125 | 42.4 | 4.267 |
| Wrist | 13.0 | 14.875 | 15.90 | 15.967 | 17.000 | 19.6 | 1.351 |
| Weight | 42.0 | 57.300 | 66.80 | 67.822 | 75.625 | 105.2 | 12.739 |

```
In [8]: #Histogram to see the frequency distribution of the variable Height  
ggplot(input_data, aes(x=Height)) +  
  geom_histogram(color="black", fill="#0002AF", alpha = 0.4, bins = 25) +  
  xlab("Height in cm") + ylab("Frequency") +  
  theme(axis.text=element_text(size=12), axis.title=element_text(size=14 , face = "bold"))
```

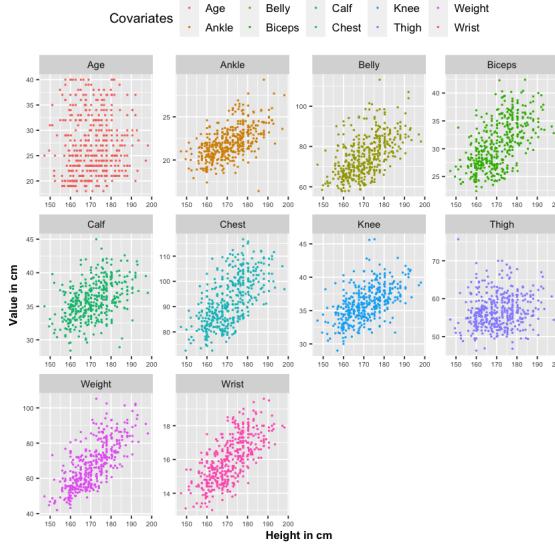


```
In [10]: corr_mat <- round(cor(input_data %>% dplyr::select(-Sex)),2)
melted_corr_mat <- melt(corr_mat) %>% dplyr::filter(Var1 == "Height")
Corr_HM <- ggplot(data = melted_corr_mat, aes(x=Var1, y=Var2,
                                              fill=value)) + geom_tile(aes(width = 1.5, height=1)) +
xlab("Response") + ylab("Covariates") + guides(fill=guide_legend(title="")) +
theme(legend.text = element_text(size=12), axis.text=element_text(size=14), axis.title=element_text(size=14 , face = "bold"))+
geom_text(aes(label = value),
          color = "white", size = 4)
Corr_HM
```



```
In [8]: #Correlation scatter plot for height vs all variables
input_corr <- input_data %>% dplyr::select(-Sex) %>%
tidyverse::gather(key = "Covariates", value = "value", -Height)

corr_plot <- ggplot(input_corr, aes(x = Height, y = value, col = Covariates)) +
  geom_point(size = 0.3) +
  facet_wrap(~Covariates, scales = "free") +
  xlab("Height in cm") + ylab("Value in cm") +
  theme(axis.text=element_text(size=6),
        axis.title=element_text(size=10 , face = "bold"),
        legend.title = element_text(size=12),
        legend.text = element_text(size=10),
        legend.position = "top",
        legend.direction = "horizontal"
      )
corr_plot
```



```
In [9]: #Linear regression model
full_model <- lm(Height ~ ., data = input_data)
summary(full_model)
```

```
Call:
lm(formula = Height ~ ., data = input_data)

Residuals:
    Min      1Q  Median      3Q     Max 
-13.9956 -2.5388  0.1734  2.7092 14.0613 

Coefficients:
            Estimate Std. Error t value Pr(>|t|)    
(Intercept) 211.23128   6.96030 30.348 < 2e-16 ***
Age          0.01769   0.03802  0.465  0.6420  
Sexm         6.38037   0.98056  6.507 2.23e-10 ***
Chest        -0.14534   0.06964 -2.087  0.0375 *  
Belly         -0.67402   0.05946 -11.335 < 2e-16 ***
Thigh        -0.58815   0.10001 -5.881 8.45e-09 ***
Knee          0.09616   0.17084  0.563  0.5738  
Calf          -0.67324   0.14570 -4.621 5.12e-06 ***
Ankle         0.10198   0.20798  0.490  0.6242  
Biceps        -1.01355   0.15004 -6.755 4.88e-11 ***
Wrist          0.52087   0.39333  1.324  0.1861  
Weight         1.40384   0.06858 20.469 < 2e-16 *** 
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 4.249 on 412 degrees of freedom
Multiple R-squared:  0.8009, Adjusted R-squared:  0.7956 
F-statistic: 150.6 on 11 and 412 DF,  p-value: < 2.2e-16
```

```
In [10]: #Finding the best subsets using least squares estimator
best_subsets <- ols_step_best_subset(full_model)
```

```
In [11]: #Listing the best subsets
best_subsets <- best_subsets %>% mutate(Model = mindex, Number_of_covariates = n,
                                             Covariates = predictors, Rsquared = rsquare, Adj.Rsquared = adjr,
                                             AIC = aic, BIC = sbc) %>%
  select(Model, Covariates, Adj.Rsquared, AIC, BIC)
best_subsets
```

A ols_step_best_subset: 11 × 5

| Model | <int> | Covariates | Adj.Rsquared | AIC | BIC |
|-------|-------|---|---|-----------|----------|
| | | <chr> | <dbl> | <dbl> | <dbl> |
| 11 | 1 | | Weight | 0.5272616 | 2789.499 |
| 51 | 2 | | Thigh Weight | 0.6630077 | 2646.977 |
| 181 | 3 | | Belly Thigh Weight | 0.7422729 | 2534.269 |
| 505 | 4 | | Belly Thigh Biceps Weight | 0.7652275 | 2495.706 |
| 841 | 5 | | Sex Belly Thigh Biceps Weight | 0.7860804 | 2457.254 |
| 1360 | 6 | | Sex Belly Thigh Calf Biceps Weight | 0.7940634 | 2442.113 |
| 1710 | 7 | | Sex Chest Belly Thigh Calf Biceps Weight | 0.7956268 | 2439.864 |
| 1949 | 8 | | Sex Chest Belly Thigh Calf Biceps Wrist Weight | 0.7965996 | 2438.820 |
| 2029 | 9 | | Sex Chest Belly Thigh Knee Calf Biceps Wrist Weight | 0.7963231 | 2440.373 |
| 2046 | 10 | | Sex Chest Belly Thigh Knee Calf Ankle Biceps Wrist Weight | 0.7959514 | 2442.121 |
| 2047 | 11 | Age Sex Chest Belly Thigh Knee Calf Ankle Biceps Wrist Weight | | 0.7955636 | 2443.898 |
| | | | | | 2496.545 |

```
In [12]: #best_subset using AIC
best_subsets %>% filter(AIC == min(best_subsets$AIC))
```

A ols_step_best_subset: 1 × 5

| Model | <int> | Covariates | Adj.Rsquared | AIC | BIC |
|-------|--|------------|--------------|---------|----------|
| | | <chr> | <dbl> | <dbl> | <dbl> |
| 8 | Sex Chest Belly Thigh Calf Biceps Wrist Weight | | 0.7965996 | 2438.82 | 2479.317 |

```
In [13]: #best_subset using BIC
best_subsets %>% filter(BIC == min(best_subsets$BIC))
```

A ols_step_best_subset: 1 × 5

| Model | <int> | Covariates | Adj.Rsquared | AIC | BIC |
|-------|------------------------------------|------------|--------------|----------|----------|
| | | <chr> | <dbl> | <dbl> | <dbl> |
| 6 | Sex Belly Thigh Calf Biceps Weight | | 0.7940634 | 2442.113 | 2474.511 |

```
In [14]: # Fitting the model, based on the based subset selection using BIC
best_subset_model <- lm(Height ~ Sex + Belly + Thigh + Calf + Biceps + Weight, data = input_data)
summary(best_subset_model)
```

Call:
`lm(formula = Height ~ Sex + Belly + Thigh + Calf + Biceps + Weight,
 data = input_data)`

Residuals:

| | | | | |
|----------|---------|--------|--------|---------|
| Min | 1Q | Median | 3Q | Max |
| -15.0336 | -2.5824 | 0.1241 | 2.6957 | 14.0164 |

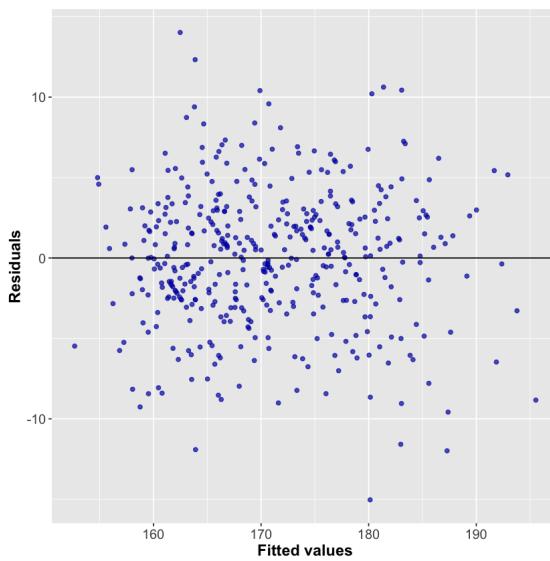
Coefficients:

| | Estimate | Std. Error | t value | Pr(> t) |
|-------------|-----------|------------|---------|--------------|
| (Intercept) | 213.72083 | 4.33769 | 49.271 | < 2e-16 *** |
| Sexm | 6.60926 | 0.94766 | 6.974 | 1.21e-11 *** |
| Belly | -0.72749 | 0.05296 | -13.737 | < 2e-16 *** |
| Thigh | -0.59827 | 0.09257 | -6.463 | 2.87e-10 *** |
| Calf | -0.52786 | 0.12726 | -4.148 | 4.07e-05 *** |
| Biceps | -1.10825 | 0.11774 | -9.413 | < 2e-16 *** |
| Weight | 1.41574 | 0.05788 | 24.461 | < 2e-16 *** |

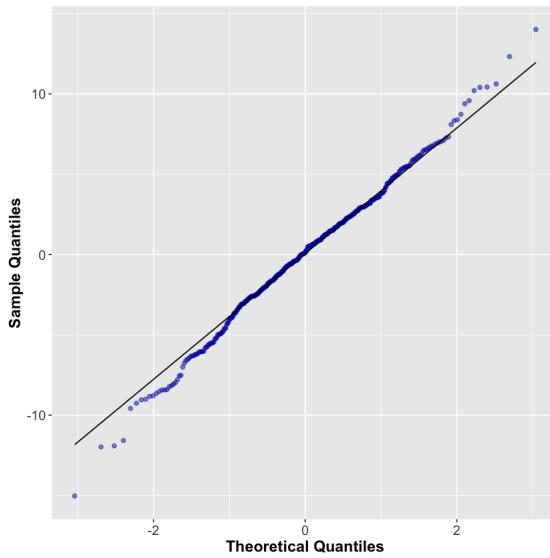
Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 4.265 on 417 degrees of freedom
 Multiple R-squared: 0.797, Adjusted R-squared: 0.7941
 F-statistic: 272.8 on 6 and 417 DF, p-value: < 2.2e-16

```
In [15]: #Residual plot for standardized residuals vs Fitted best subset model
ggplot(best_subset_model, aes(x = .fitted, y = best_subset_model$residuals)) +
  geom_point(size = 1.5, color = "#0002AF", alpha = 0.7) +
  geom_hline(yintercept = mean(best_subset_model$residuals), color = "black", linewidth = 0.5) + xlab("Fitted values") + ylab("R
    theme(axis.text=element_text(size=12), axis.title=element_text(size=14 , face = "bold"))
```



```
In [16]: # QQ plot for the standardized residuals
ggplot(input_data, aes(sample = best_subset_model$residuals)) +
  stat_qq(distribution = stats::qnorm, color = "#0002AF", alpha = 0.5) +
  stat_qq_line() +
  xlab("Theoretical Quantiles") +
  ylab("Sample Quantiles") +
  theme(plot.title = element_text(face = "bold", hjust = 0.5, size = 12),
        axis.text=element_text(size=12),
        axis.title=element_text(size=14, face = "bold"),
        legend.text = element_text(size = 12))
```



```
In [17]: #Extracting the coefficient values and p-values
coefficient_values <- t(rbind(round(best_subset_model$coefficients,3)))
p_values <- t(rbind(round(summary(best_subset_model)$coefficients[, 4],10)))
```

```
In [18]: #Finding the confidence intervals for the estimated coefficients
conf_intervals <- as.data.frame(round(confint(best_subset_model, level = 0.95),3))
colnames(conf_intervals) <- c("lower","upper")
conf_intervals['Confidence_interval'] <- paste0("[",conf_intervals$lower, " , ", conf_intervals$upper, "]")
conf_intervals <- conf_intervals %>% select(Confidence_interval)
```

```
In [19]: #Displaying the output
```

```
parameters <- cbind(coefficient_values, conf_intervals, p_values)
colnames(parameters) <- c("Estimated coefficients", "Confidence intervals", "P_values")
parameters
```

A data.frame: 7 × 3

| | Estimated coefficients | Confidence intervals | P_values |
|-------------|------------------------|----------------------|-------------|
| | <dbl> | <chr> | <dbl> |
| (Intercept) | 213.721 | [205.194 , 222.247] | 0.00000e+00 |
| Sexm | 6.609 | [4.746 , 8.472] | 0.00000e+00 |
| Belly | -0.727 | [-0.832 , -0.623] | 0.00000e+00 |
| Thigh | -0.598 | [-0.78 , -0.416] | 3.00000e-10 |
| Calf | -0.528 | [-0.778 , -0.278] | 4.07118e-05 |
| Biceps | -1.108 | [-1.34 , -0.877] | 0.00000e+00 |
| Weight | 1.416 | [1.302 , 1.53] | 0.00000e+00 |